

Programming for Evolutionary Biology
March 21 – April 6 2014
Leipzig, Germany

Introduction to Unix systems

Part 4 – awk & make

Giovanni Marco Dall'Olio
King's College of London, UK

Schedule

- 9.30 – 10.30: “What is Unix?”
- 11.00 – 12.00: Introducing the terminal
- 13.30 – 16.30: Grep, & Unix philosophy
- **17:00 – 18:00: awk, make, and question time**

awk

- **awk** is the Unix command to work with tabular data

awk

- Things you can do with awk:
 - Extract all the lines of a file that match a pattern, and print only some of the columns (instead of “grep | cut”)
 - Add a prefix/suffix to all the element of a column (instead of “cut | paste”)
 - Sum values of different columns

awk and gawk

- Today we will use the GNU version of **awk**, called **gawk**
- In most systems, awk and gawk are the same

Awk exercise

- Let's go to the “bed” directory
- `cd .homes/evopserver/lectures/unix_intro/bed`

The BED format

- The BED format is used to store annotations on genomic coordinates
- Example:
 - Annotate gene position
 - Annotate Transcription Factor binding sites
 - Annotate SNP genotypes
 - Annotate Gene Expression

Example awk usage

- Select only the lines matching chromosome 7, and print only the second column, summing 100 to it
 - `awk '$1 ~ chr7 {print $2+100}' annotations.bed`

Example of BED file

Columns in a BED file:

- **Chromosome**
- **Start**
- **End**
- **Label**
- **Score**
- **Strand (+ or -)**
-

chr7	127471196	127472363	gene1	0	+
chr7	127472363	127473530	gene2	2	+
chr7	127473530	127474697	gene3	20	+
chr7	127474697	127475864	gene4	3	+
chr7	127475864	127477031	gene5	100	-
chr7	127477031	127478198	gene6	3	-
chr7	127478198	127479365	gene7	5	-
chr7	127479365	127480532	gene8	1	+
chr7	127480532	127481699	gene9	3	-

Basic awk usage

- `awk '<pattern to select lines> {instructions to be executed on each line}'`

Selecting columns in awk

- In awk, each column can be accessed by `$<column-number>`
- For example,
 - `$1` → first column of the file
 - `$2` → second column
 - `$NF` → last column
- `$0` matches all the columns of the file

Printing columns in awk

- The basic awk usage is to print specific columns
- The syntax is the following:
 - `awk '{print $1, $2, $3}' annotations.bed` → prints the first three columns

Printing columns in awk

- The basic awk usage is to print specific columns
- The syntax is the following:
 - `awk '{print $1, $2, $3}' annotations.bed` → prints the first three columns
 - `awk '{print $0}' annotations.bed` → print all the columns

Adding a prefix to a column with awk

- A common awk usage is to add a prefix or suffix to all the entries of a column
- Example:
 - `awk '{print $2 "my_prefix"$2}' annotations.bed`

Summing columns in awk

- If two columns contain numeric values, we can use awk to sum them
- Usage:
 - `awk '{print $2 + $3}' annotations.bed` → sums columns 2 and 3

Filter lines and select columns with awk

- Awk can apply a filter and print only the lines matching a pattern
 - Like grep, but more powerful
- Print all the lines that have “chr7” in their first column:
 - `awk '$1 ~ “chr7” {print $0}' myfile.txt`

Advanced: redirecting to files

- The following will split the contents of annotations.bed into two files, according to the 1st column:
- `awk '{print $0 > $1".file"}' annotations.bed`

Advanced: print something before reading the file

- The BEGIN statement can be used to execute commands before reading the file:
- `awk 'BEGIN {print "position"} {print $2}' annotations.bed`

More on awk

- awk is a complete programming language
- It is the equivalent of a spreadsheet for the command line
- If you want to know more, check the book “Gawk effective AWK Programming” at <http://www.gnu.org/software/gawk/manual>

A streaming file editor: sed

- **sed** is a command-line tool to edit files, without opening their full contents in memory
- Things you can achieve with sed:
 - Find&Replace words in huge text files
 - Remove determinate lines from a file

sed - basic usage

- The following will replace all the occurrences of the word 'gene' with 'GENE' in the BED file:
 - `sed 's/gene/GENE/' annotations.bed`
- Explanation:
 - `sed` → name of the command
 - `'s/gene/GENE/'` → substitute (s) the word 'gene' (first pair of slashes) with the word 'GENE' (second pair of slashes)

Sed - saving results to file

- The previous command printed the result of the substitution, but did not save it to a file
- If you open the file annotations.bed, it has not been changed
- To make your changes permanents, you have two options:
 - Output the result to a file:
`sed 's/gene/GENE/' annotations.bed > annotations_changed.bed`
 - Use the -i option:
`sed -i 's/gene/GENE/' annotations.bed`

Other sed options: removing a pattern of lines

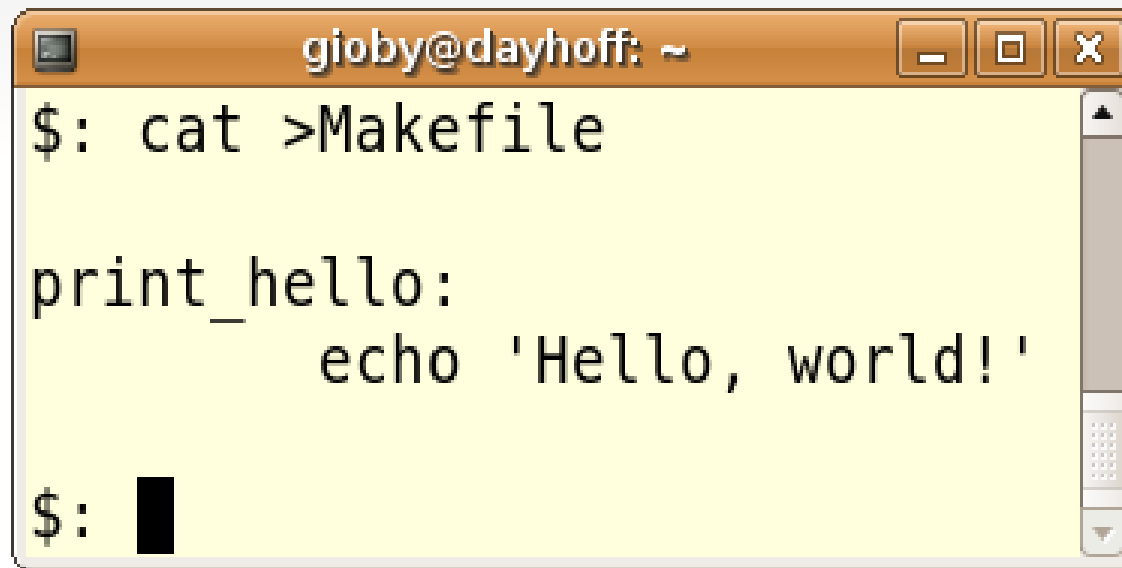
- Let's remove all the 1,3,5,7, etc.. lines from the bed file:
 - `sed -n '1~2d' annotations.bed`

GNU/make

- make is a tool used by programmer to define how software should be compiled and installed
- It is also used as a way to store a set of commands, to recall them later

Simplest Makefile example

- The simplest Makefile contains just the name of a task and the commands associated with it:

A terminal window with a title bar that reads "gioby@dayhoff: ~". The window has standard minimize, maximize, and close buttons. The terminal text is as follows:

```
$: cat >Makefile  
  
print_hello:  
    echo 'Hello, world!'  
  
$: █
```

- `print_hello` is a makefile 'rule': it stores the commands needed to say 'Hello, world!' to the screen.

Simplest Makefile example



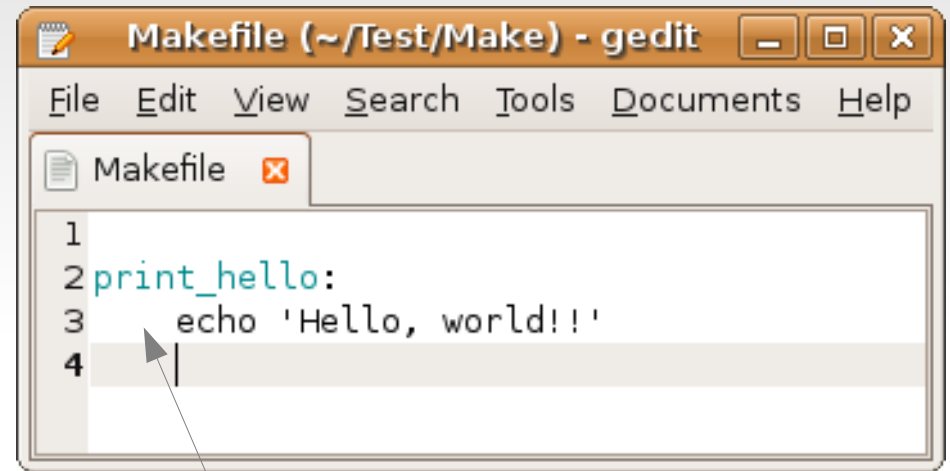
Simplest Makefile example

- Create a file in your computer and save it as 'Makefile'.
- Write these instructions in it:

```
print_hello:  
    echo 'Hello, world!!'
```

- Then, open a terminal and type:

make -f Makefile print_hello



This is a tabulation
(<Tab> key)

Simplest Makefile example

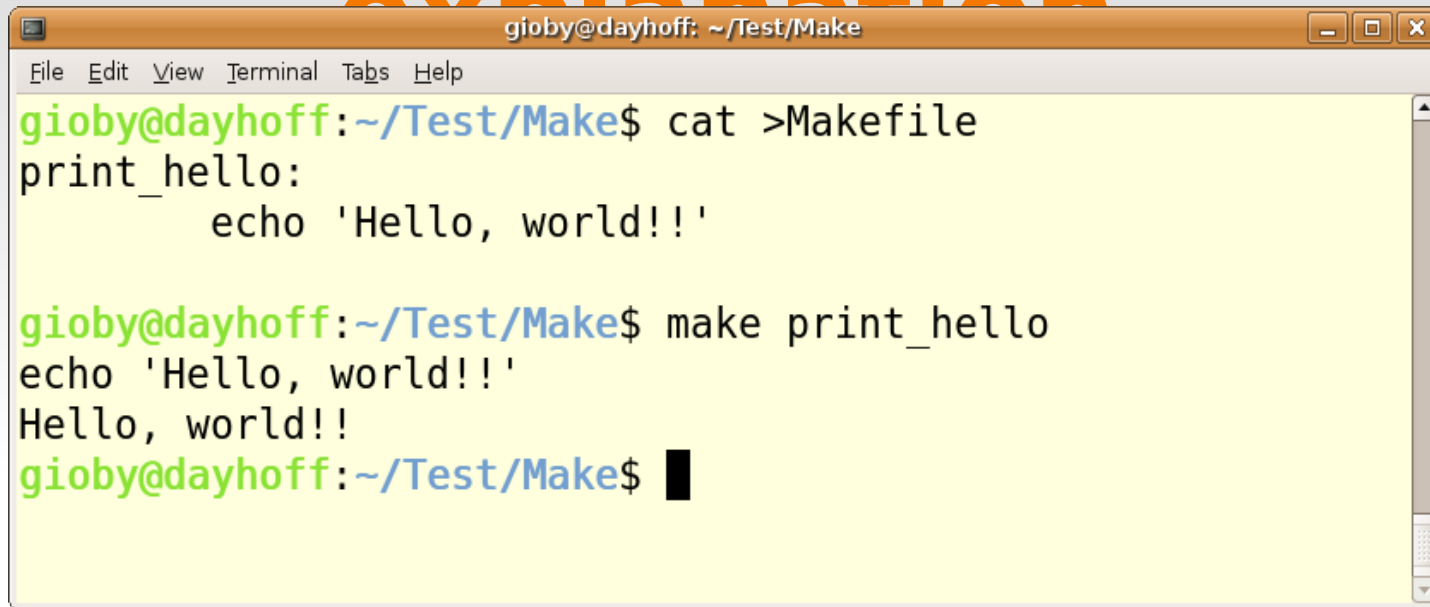
A terminal window titled 'gioby@dayhoff: ~' with standard window controls. The terminal shows the creation of a Makefile, its execution, and the resulting output.

```
File Edit View Terminal Tabs Help
$: cat >Makefile

print_hello:
    echo "Hello, world!!"

$: make -f Makefile print_hello
echo "Hello, world!!"
Hello, world!!
$:
```

Simplest Makefile example - explanation

A terminal window titled 'gioby@dayhoff: ~/Test/Make' with a menu bar (File, Edit, View, Terminal, Tabs, Help). The terminal shows the following commands and output:

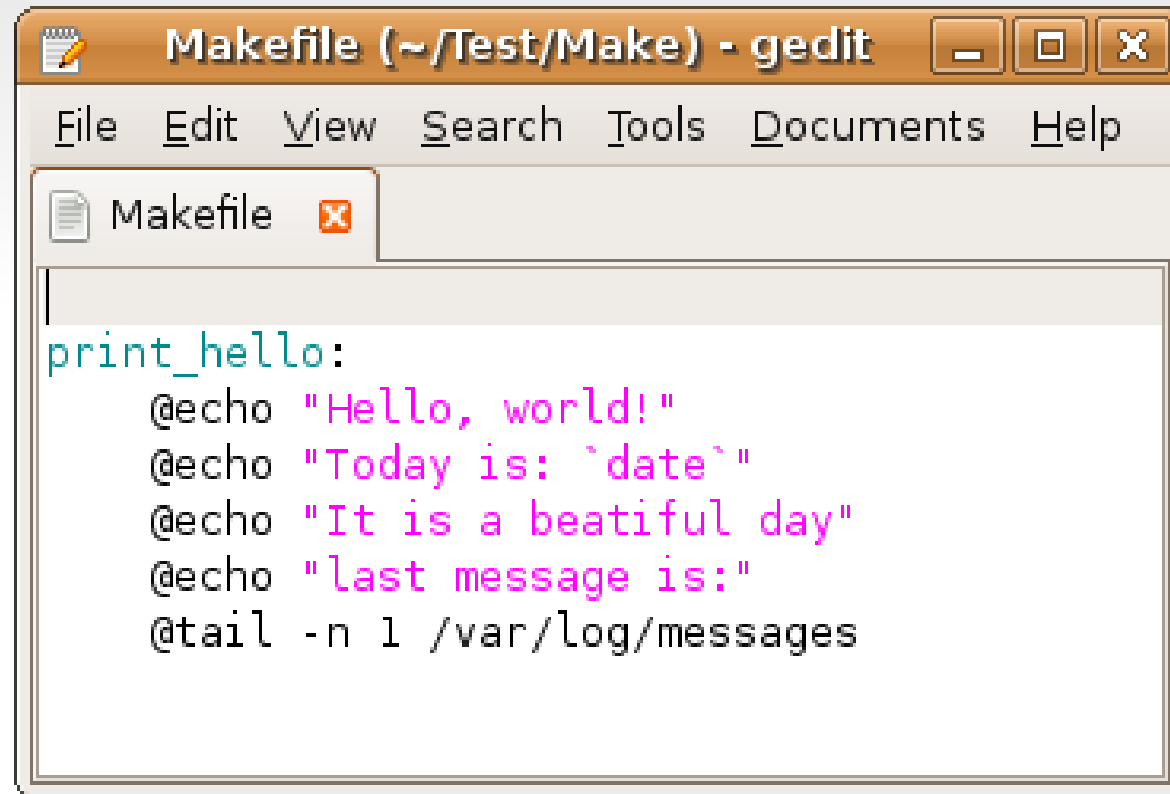
```
gioby@dayhoff:~/Test/Make$ cat >Makefile
print_hello:
    echo 'Hello, world!!'

gioby@dayhoff:~/Test/Make$ make print_hello
echo 'Hello, world!!'
Hello, world!!
gioby@dayhoff:~/Test/Make$
```

- When invoked, the program 'make' looks for a file in the current directory called 'Makefile'
- When we type 'make print_hello', it executes any procedure (target) called 'print_hello' in the makefile
- It then shows the commands executed and their output

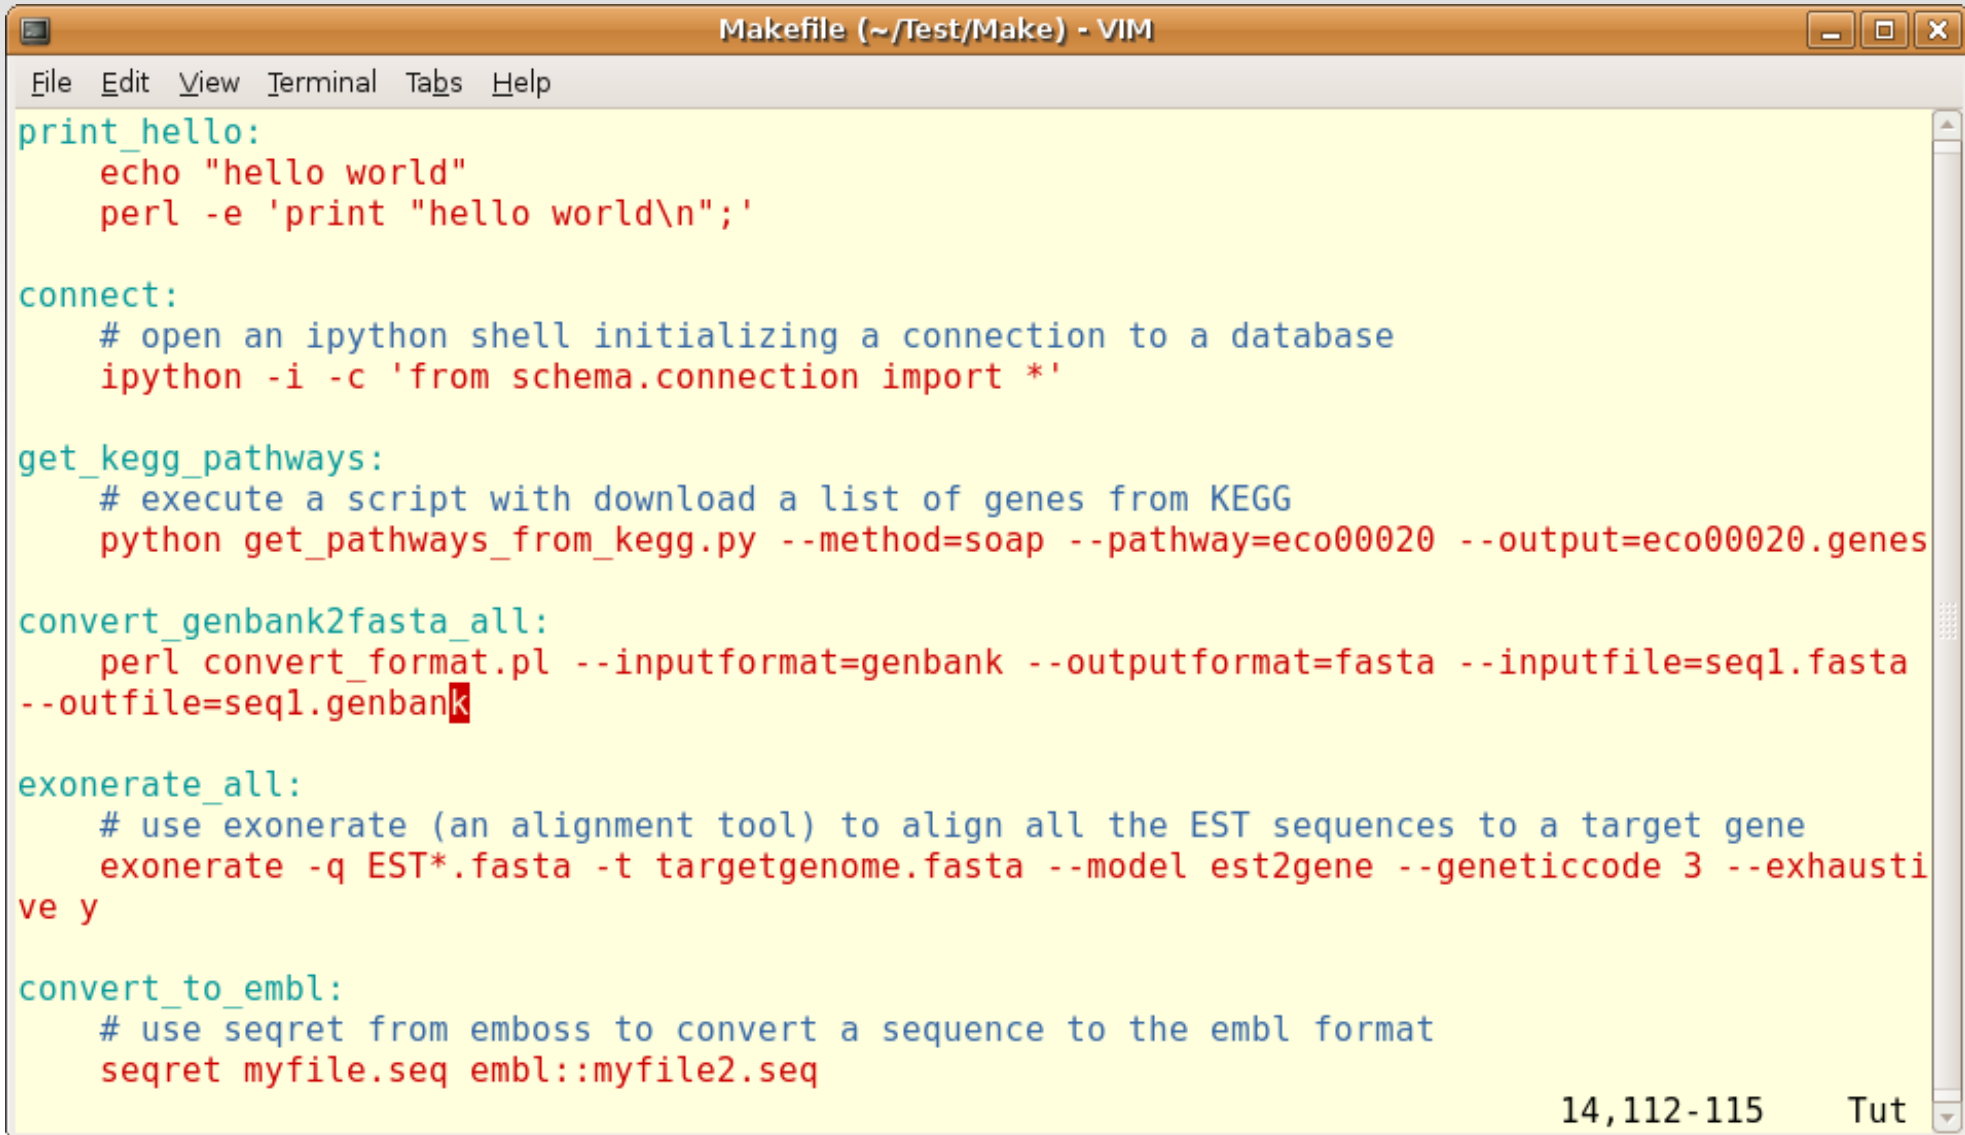
A slightly longer example

- You can add as many commands you like to a rule
- For example, this 'print_hello' rule contains 5 commands
- Note: ignore the '@' thing, it is only to disable verbose mode (explained later)

A screenshot of a gedit text editor window titled "Makefile (~/.Test/Make) - gedit". The window has a menu bar with "File", "Edit", "View", "Search", "Tools", "Documents", and "Help". Below the menu bar is a tab labeled "Makefile" with a red 'x' icon. The main text area contains the following code:

```
print_hello:
    @echo "Hello, world!"
    @echo "Today is: `date`"
    @echo "It is a beautiful day"
    @echo "last message is:"
    @tail -n 1 /var/log/messages
```

A more complex example



```
Makefile (~/.Test/Make) - VIM
File Edit View Terminal Tabs Help

print_hello:
    echo "hello world"
    perl -e 'print "hello world\n";'

connect:
    # open an ipython shell initializing a connection to a database
    ipython -i -c 'from schema.connection import *'

get_kegg_pathways:
    # execute a script with download a list of genes from KEGG
    python get_pathways_from_kegg.py --method=soap --pathway=eco00020 --output=eco00020.genes

convert_genbank2fasta_all:
    perl convert_format.pl --inputformat=genbank --outputformat=fasta --inputfile=seq1.fasta
    --outfile=seq1.genbank

exonerate_all:
    # use exonerate (an alignment tool) to align all the EST sequences to a target gene
    exonerate -q EST*.fasta -t targetgenome.fasta --model est2gene --geneticcode 3 --exhaustive y

convert_to_embl:
    # use seqret from emboss to convert a sequence to the embl format
    seqret myfile.seq embl::myfile2.seq
```

14,112-115 Tut

The target syntax

- The target of a rule can be either a title for the task, or a file name.
- Every time you call a make rule (example: 'make all'), the program looks for a file called like the target name (e.g. 'all', 'clean', 'inputdata.txt', 'results.txt')
- The rule is executed only if that file doesn't exist.

Filename as target names



```
gioby@dayhoff: ~  
$: cat >Makefile  
  
testfile.txt:  
    @touch testfile.txt  
    @echo 'testfile.txt has been created'  
  
clean:  
    @rm testfile.txt  
    @echo 'testfile.txt has been deleted'  
  
$: █
```

- In this makefile, we have two rules: 'testfile.txt' and 'clean'

Filename as target names

A terminal window with a yellow background and a brown title bar. The title bar contains the text 'gioby@dayhoff: ~' and standard window control buttons. The terminal shows the command '\$: cat >Makefile' followed by the contents of a Makefile. The Makefile has two rules: 'testfile.txt' and 'clean'. The 'testfile.txt' rule has two commands: '@touch testfile.txt' and '@echo 'testfile.txt has been created''. The 'clean' rule has two commands: '@rm testfile.txt' and '@echo 'testfile.txt has been deleted''. The prompt '\$: ' is followed by a black cursor block.

```
gioby@dayhoff: ~
$: cat >Makefile

testfile.txt:
    @touch testfile.txt
    @echo 'testfile.txt has been created'

clean:
    @rm testfile.txt
    @echo 'testfile.txt has been deleted'

$: █
```

- In this makefile, we have two rules: 'testfile.txt' and 'clean'
- When we call 'make testfile.txt', make checks if a file called 'testfile.txt' already exists.

Filename as target names

```
gioby@dayhoff: ~  
$: cat >Makefile  
  
testfile.txt:  
    @touch testfile.txt  
    @echo 'testfile.txt has been created'  
  
clean:  
    @rm testfile.txt  
    @echo 'testfile.txt has been deleted'  
  
$: make testfile.txt  
testfile.txt has been created  
$: make testfile.txt  
make: `testfile.txt' is up to date.  
$: █
```

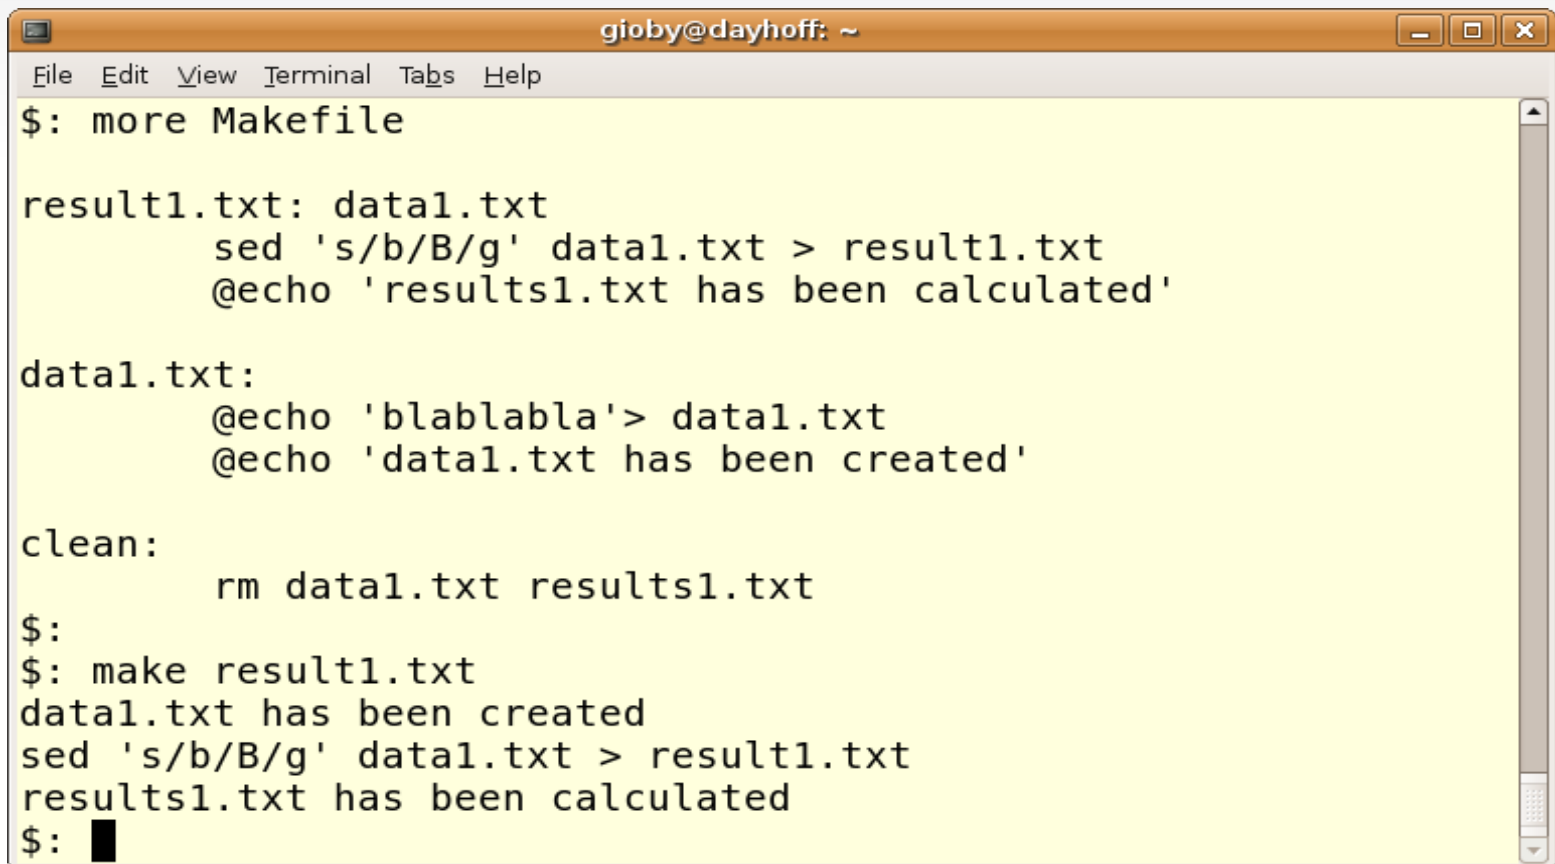
The commands associated with the rule '**testfile.txt**' are executed only if that file doesn't exist already

Real Makefile-rule syntax

- Complete syntax for a Makefile rule:
 <target>: <list of prerequisites>
 <commands associated to the rule>
- Example:
 result1.txt: data1.txt data2.txt
 cat data1.txt data2.txt > result1.txt
 @echo 'result1.txt' has been calculated'
- Prerequisites are files (or rules) that need to exist already in order to create the target file.
- If 'data1.txt' and 'data2.txt' don't exist, the rule 'result1.txt' will exit with an error (no rule to create them)

Piping Makefile rules together

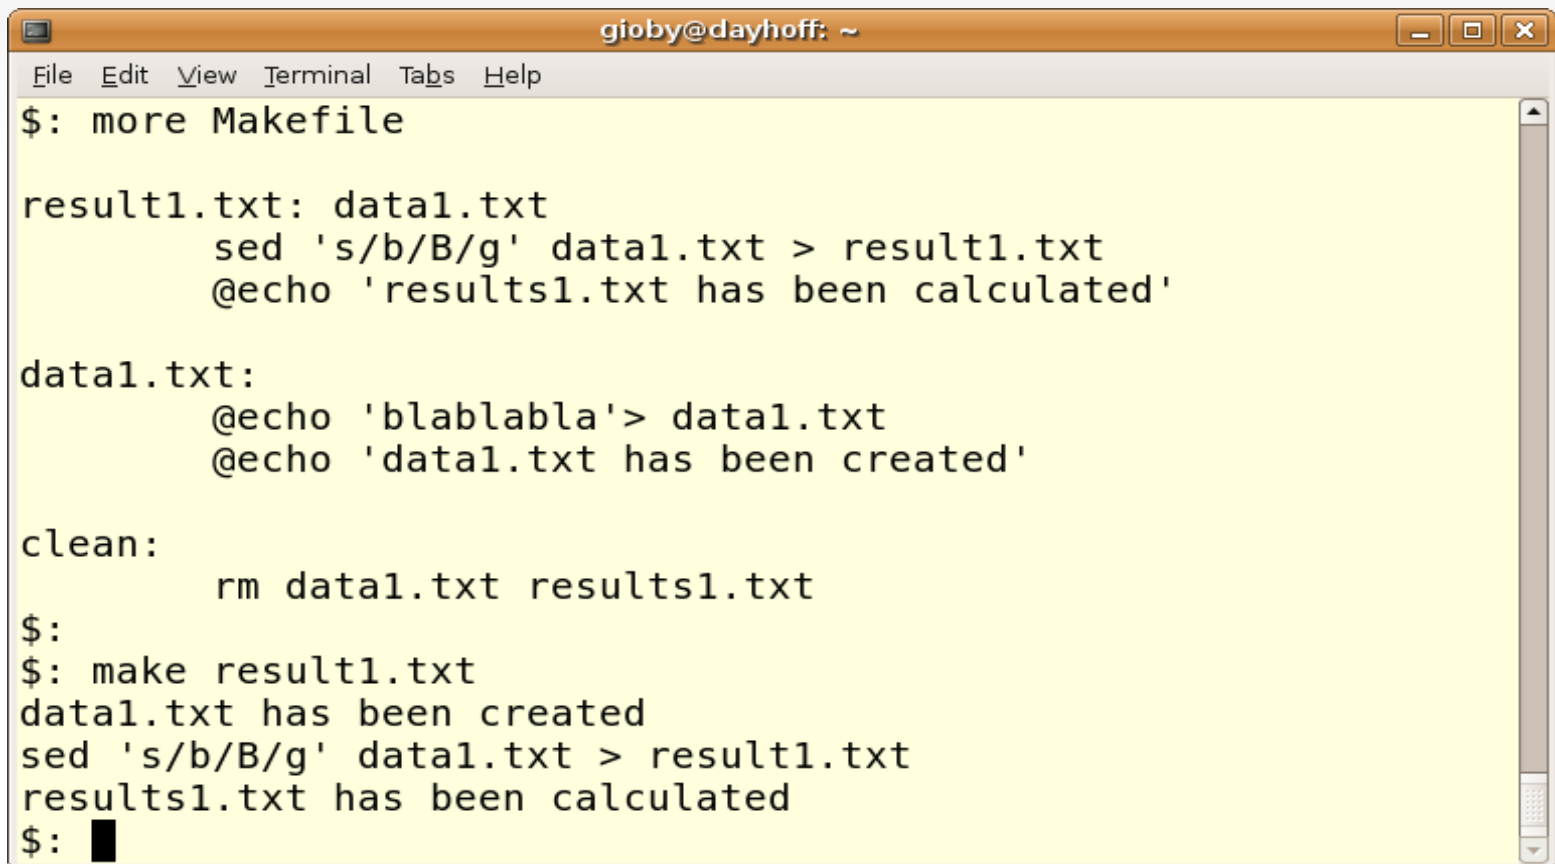
- You can pipe two Makefile rules together by defining prerequisites



```
gioby@dayhoff: ~  
File Edit View Terminal Tabs Help  
$: more Makefile  
  
result1.txt: data1.txt  
    sed 's/b/B/g' data1.txt > result1.txt  
    @echo 'results1.txt has been calculated'  
  
data1.txt:  
    @echo 'blablabla'> data1.txt  
    @echo 'data1.txt has been created'  
  
clean:  
    rm data1.txt results1.txt  
  
$:  
$: make result1.txt  
data1.txt has been created  
sed 's/b/B/g' data1.txt > result1.txt  
results1.txt has been calculated  
$:
```

Piping Makefile rules together

- The rule 'result1.txt' depends on the rule 'data1.txt', which should be executed first



```
gioby@dayhoff: ~  
File Edit View Terminal Tabs Help  
$: more Makefile  
  
result1.txt: data1.txt  
    sed 's/b/B/g' data1.txt > result1.txt  
    @echo 'results1.txt has been calculated'  
  
data1.txt:  
    @echo 'blablabla'> data1.txt  
    @echo 'data1.txt has been created'  
  
clean:  
    rm data1.txt results1.txt  
  
$:  
$: make result1.txt  
data1.txt has been created  
sed 's/b/B/g' data1.txt > result1.txt  
results1.txt has been calculated  
$: █
```

Conditional execution by modification date

- We have seen how make can be used to create a file, if it doesn't exist.

file.txt:

```
# if file.txt doesn't exist, then create it:  
echo 'contents of file.txt' > file.txt
```

- We can do better: create or update a file only if it is newer than its prerequisites

Conditional execution by modification date

- Let's have a better look at this example:

```
result1.txt: data1.txt calculate_result.py  
python calculate_result.txt --input  
data1.txt
```

- A great feature of make is that it execute a rule not only if the target file doesn't exist, but also if it has a 'last modification date' earlier than all of its prerequisites

Conditional execution by modification date

```
result1.txt: data1.txt  
    @sed 's/b/B/i' data1.txt > result1.txt  
    @echo 'result1.txt has been calculated'
```

- In this example, result1.txt will be recalculated every time 'data1.txt' is modified

```
$: touch data1.txt calculate_result.py
```

```
$: make result1.txt  
result1.txt has been calculated
```

```
$: make result1.txt  
result1.txt is already up-to-date
```

```
$: touch data1.txt  
$: make result1.txt  
result1.txt has been calculated
```

Make - advantages

- Make allows you to save shell commands along with their parameters and re-execute them;
- It allows you to use command-line tools which are more flexible;
- Combined with a revision control software, it makes possible to reproduce all the operations made to your data;

Resume of the day

- Unix → an operating system from the '70s, which was successful for its philosophy and technical novelties
- The terminal → a way to execute commands by typing
- How to get help → man, info, --help, internet
- Grep, awk → search patterns in a file
- Other Unix tools → each specialized on a single data manipulation task